Unsupervised Learning: Word Embedding

1-of-N Encoding

apple = [1 0 0 0 0]

bag = [0 1 0 0 0]

cat = [0 0 1 0 0]

dog = [0 0 0 1 0]

elephant = [0 0 0 0 1]

Word Embedding

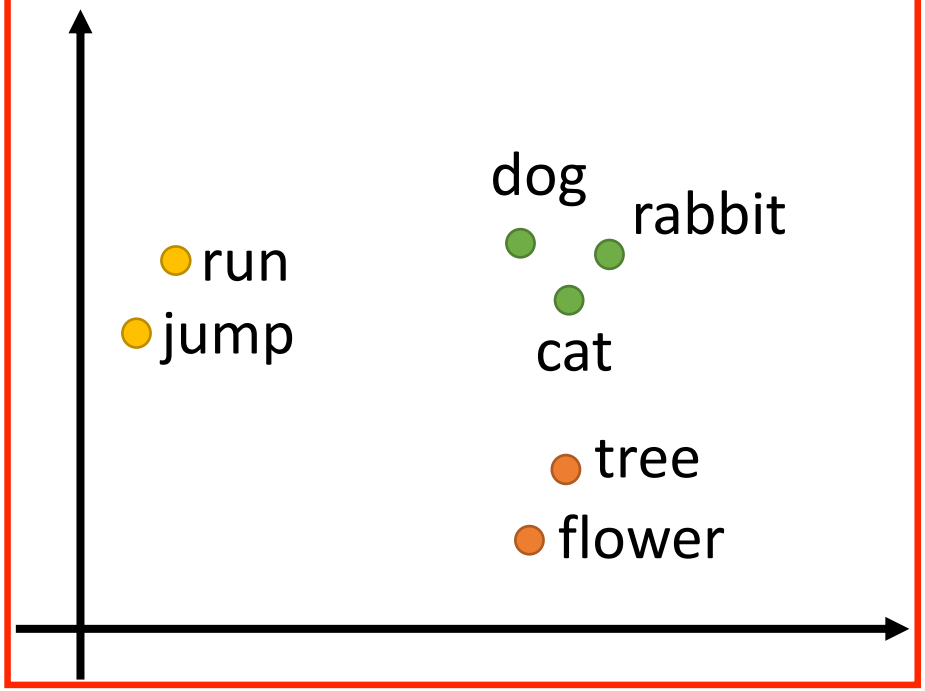

Word Class

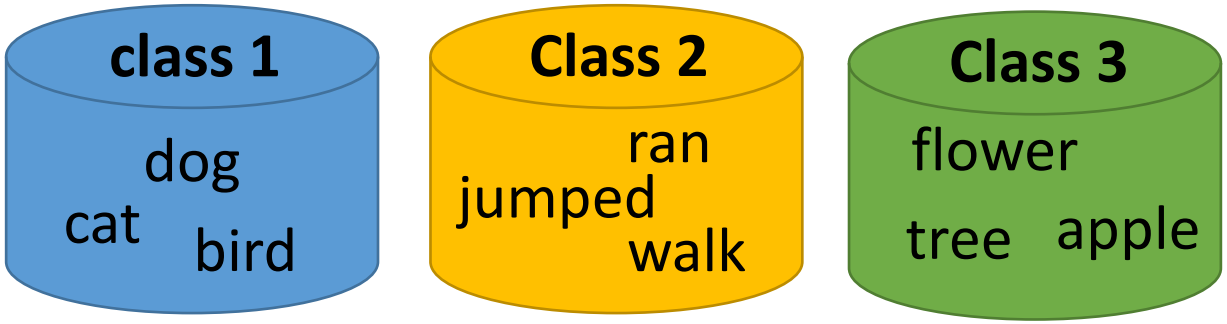

Word Embedding

- Machine learns the meaning of words from reading a lot of documents without supervision

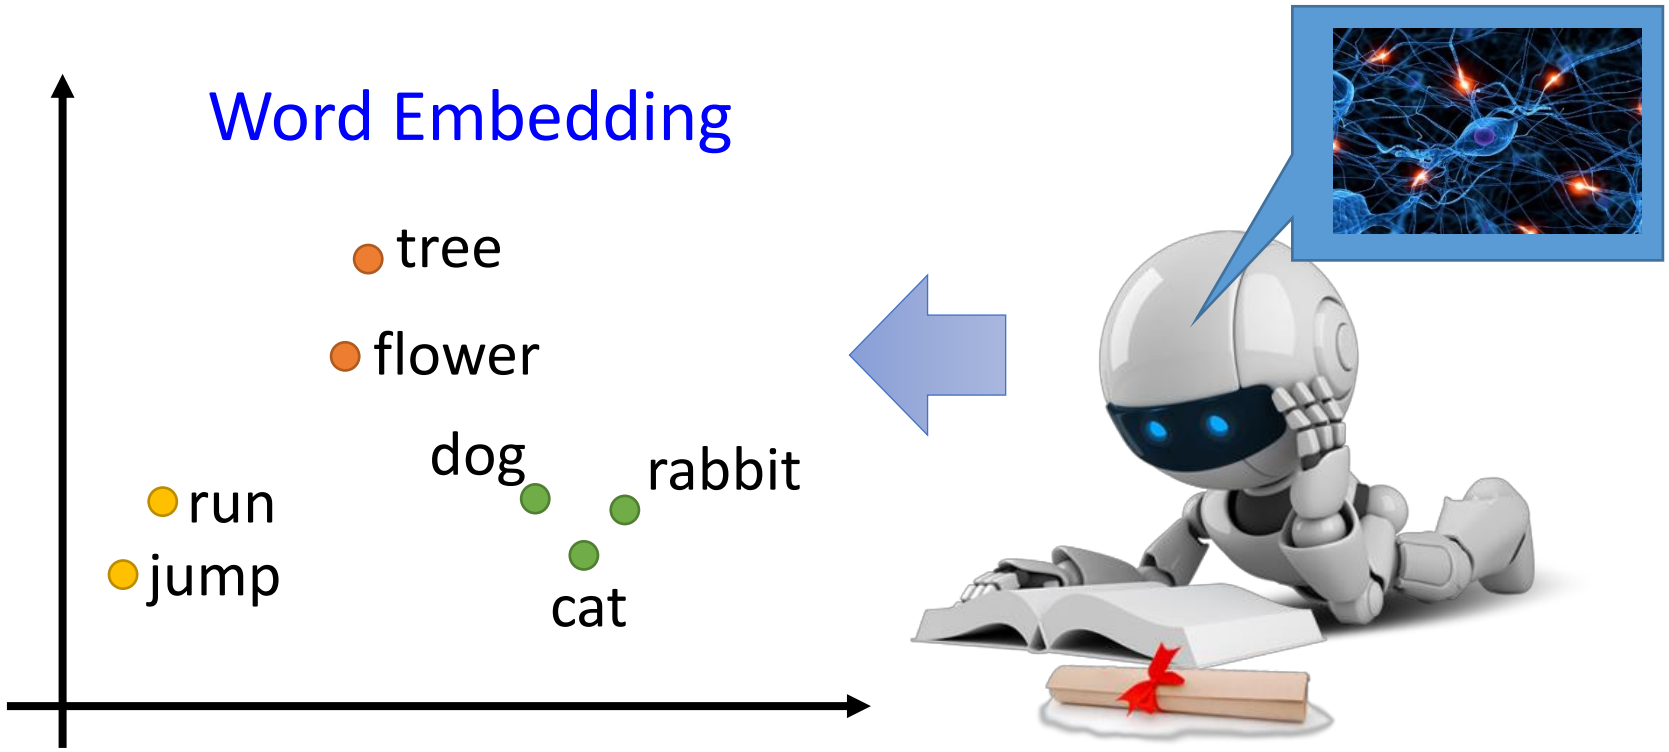

Word Embedding

- Machine learns the meaning of words from reading a lot of documents without supervision
- A word can be understood by its context

蔡英文、馬英九 are something very similar

You shall know a word by the company it keeps

馬英九 520宣誓就職

蔡英文 520宣誓就職

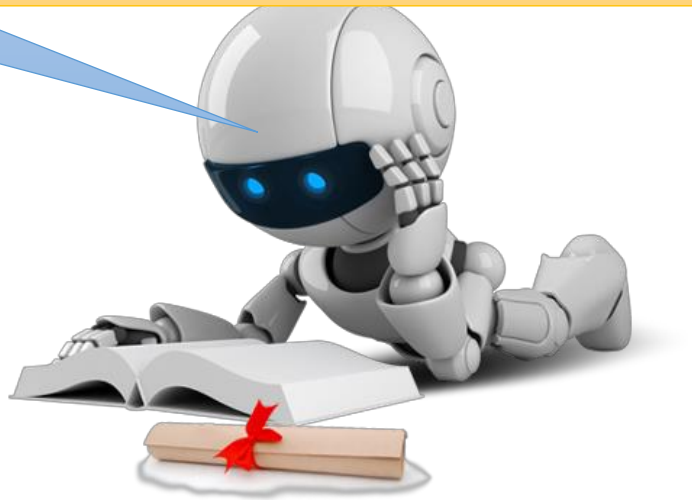

How to exploit the context?

進行 Word Embedding（尋找一個 Word 的 Vector）有兩種方式：

- **Count based**
  - If two words w_i and w_j frequently co-occur, $V(w_i)$ and $V(w_j)$ would be close to each other
  - E.g. Glove Vector: http://nlp.stanford.edu/projects/glove/

$V(w_i) \cdot V(w_j)$ ⟷ $N_{i,j}$

Inner product

Number of times w_i and w_j in the same document

- **Prediction based**

Prediction-based

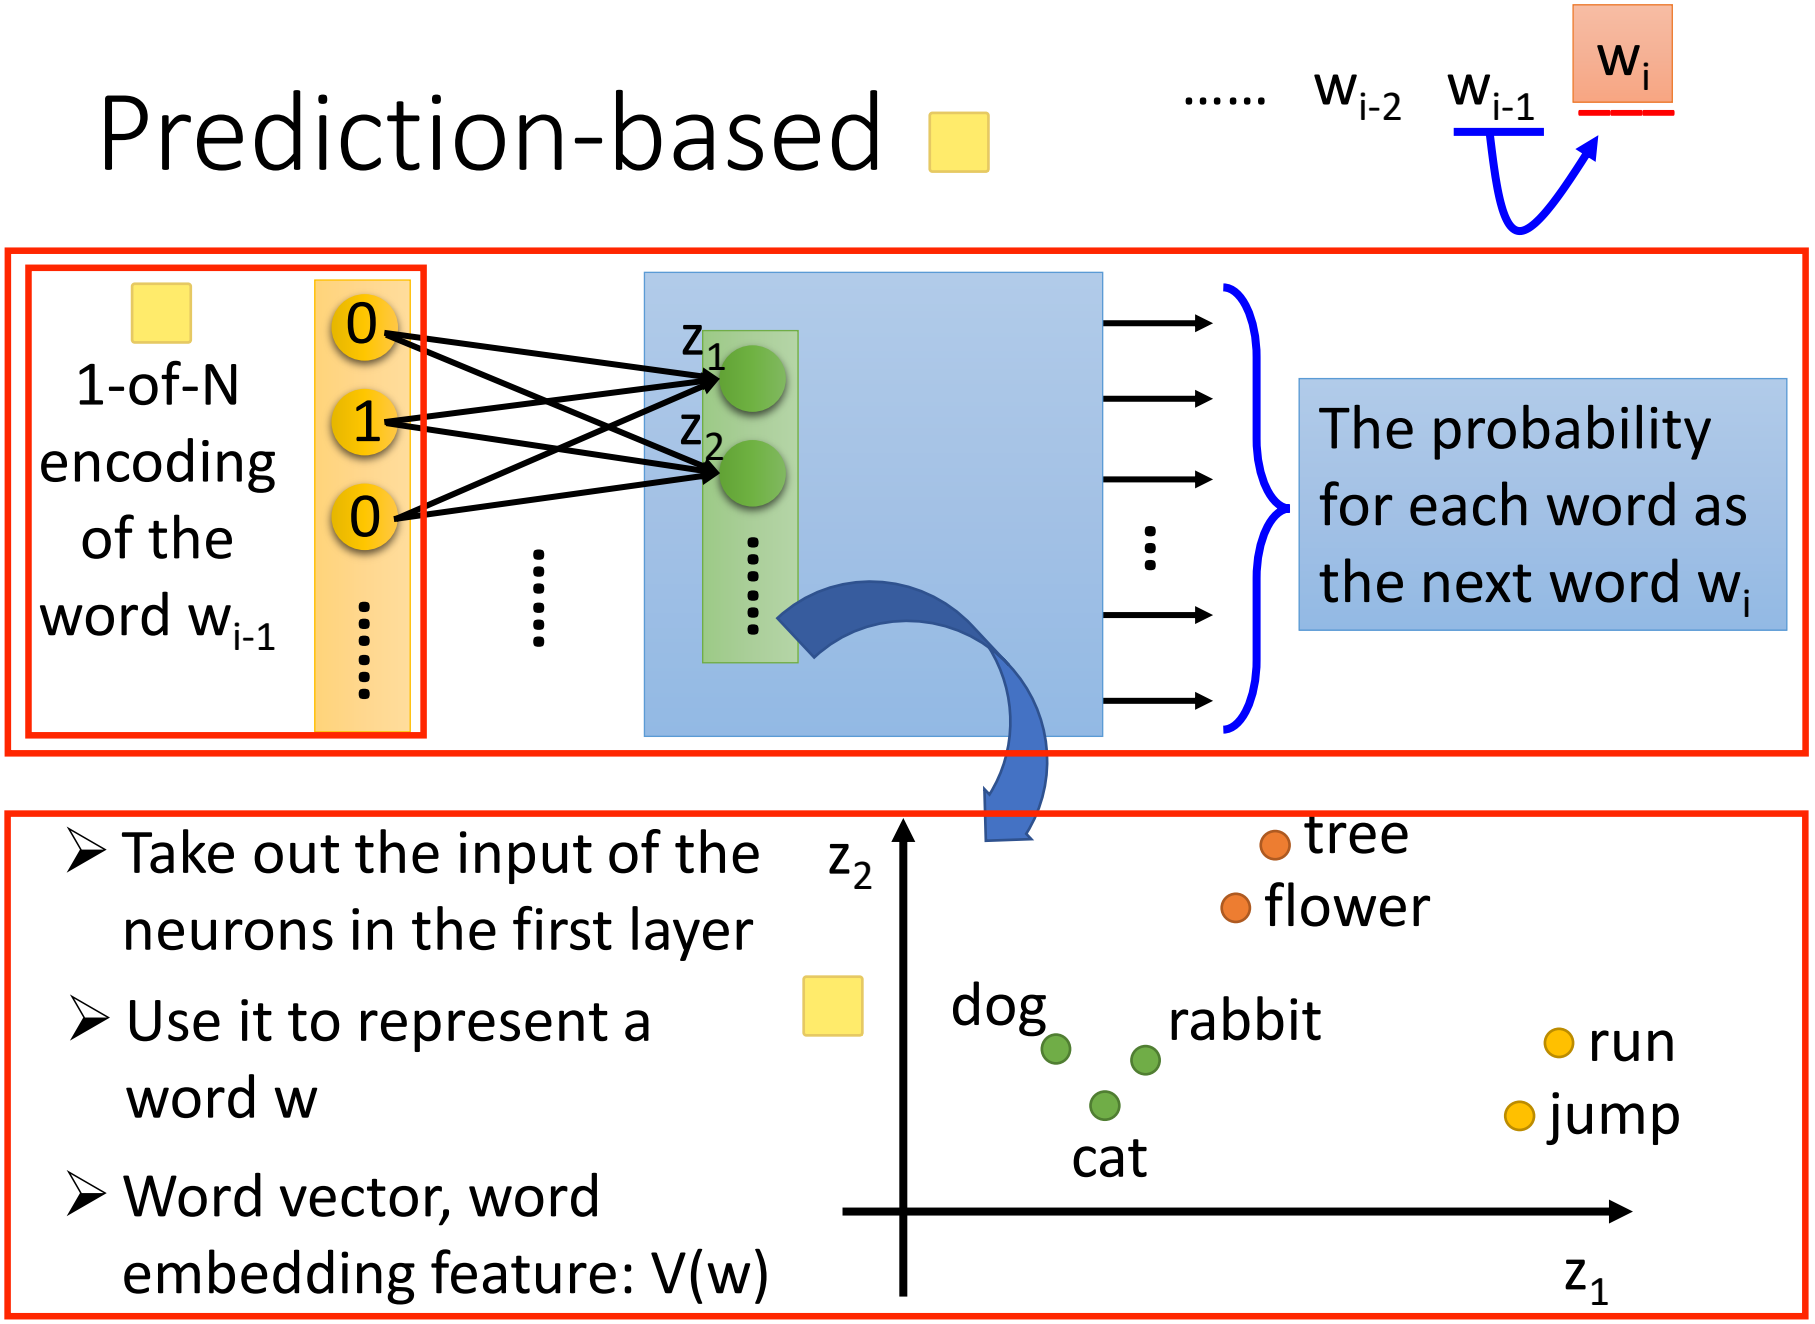

Prediction-based

You shall know a word by the company it keeps

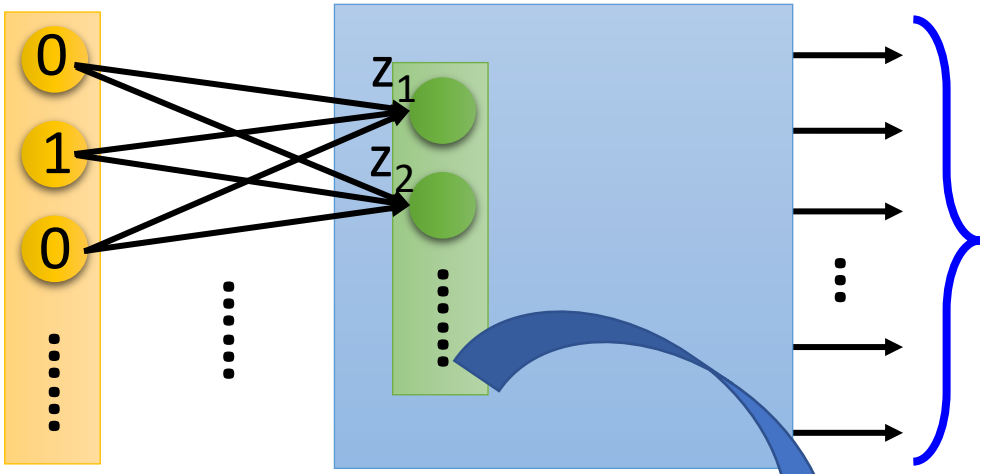

蔡英文
or
馬英九

The probability for each word as the next word w_i

"宣誓就職" should have large probability

Training text:

..... 蔡英文 宣誓就職 .....
w_{i-1} w_i

..... 馬英九 宣誓就職 .....
w_{i-1} w_i

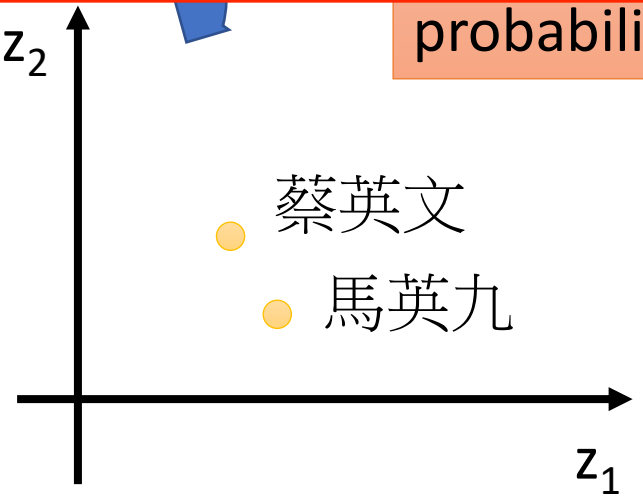

Prediction-based – Sharing Parameters

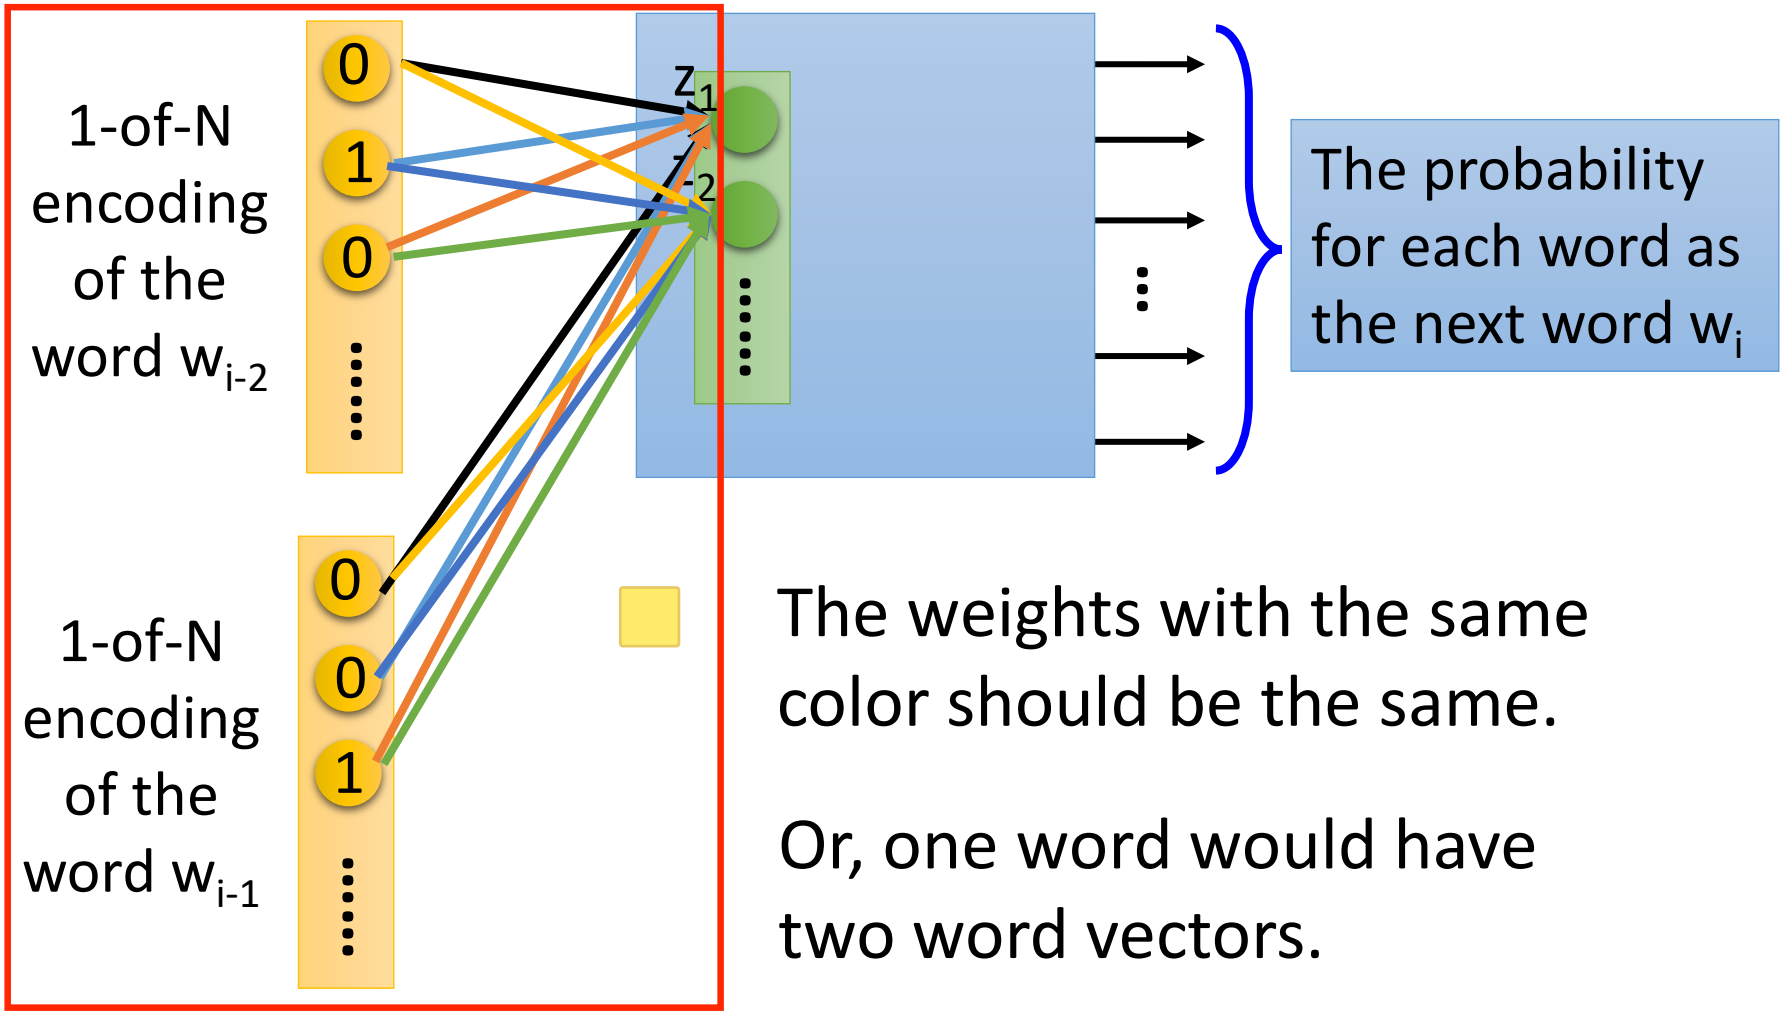

The weights with the same color should be the same.

Or, one word would have two word vectors.

Prediction-based – Sharing Parameters

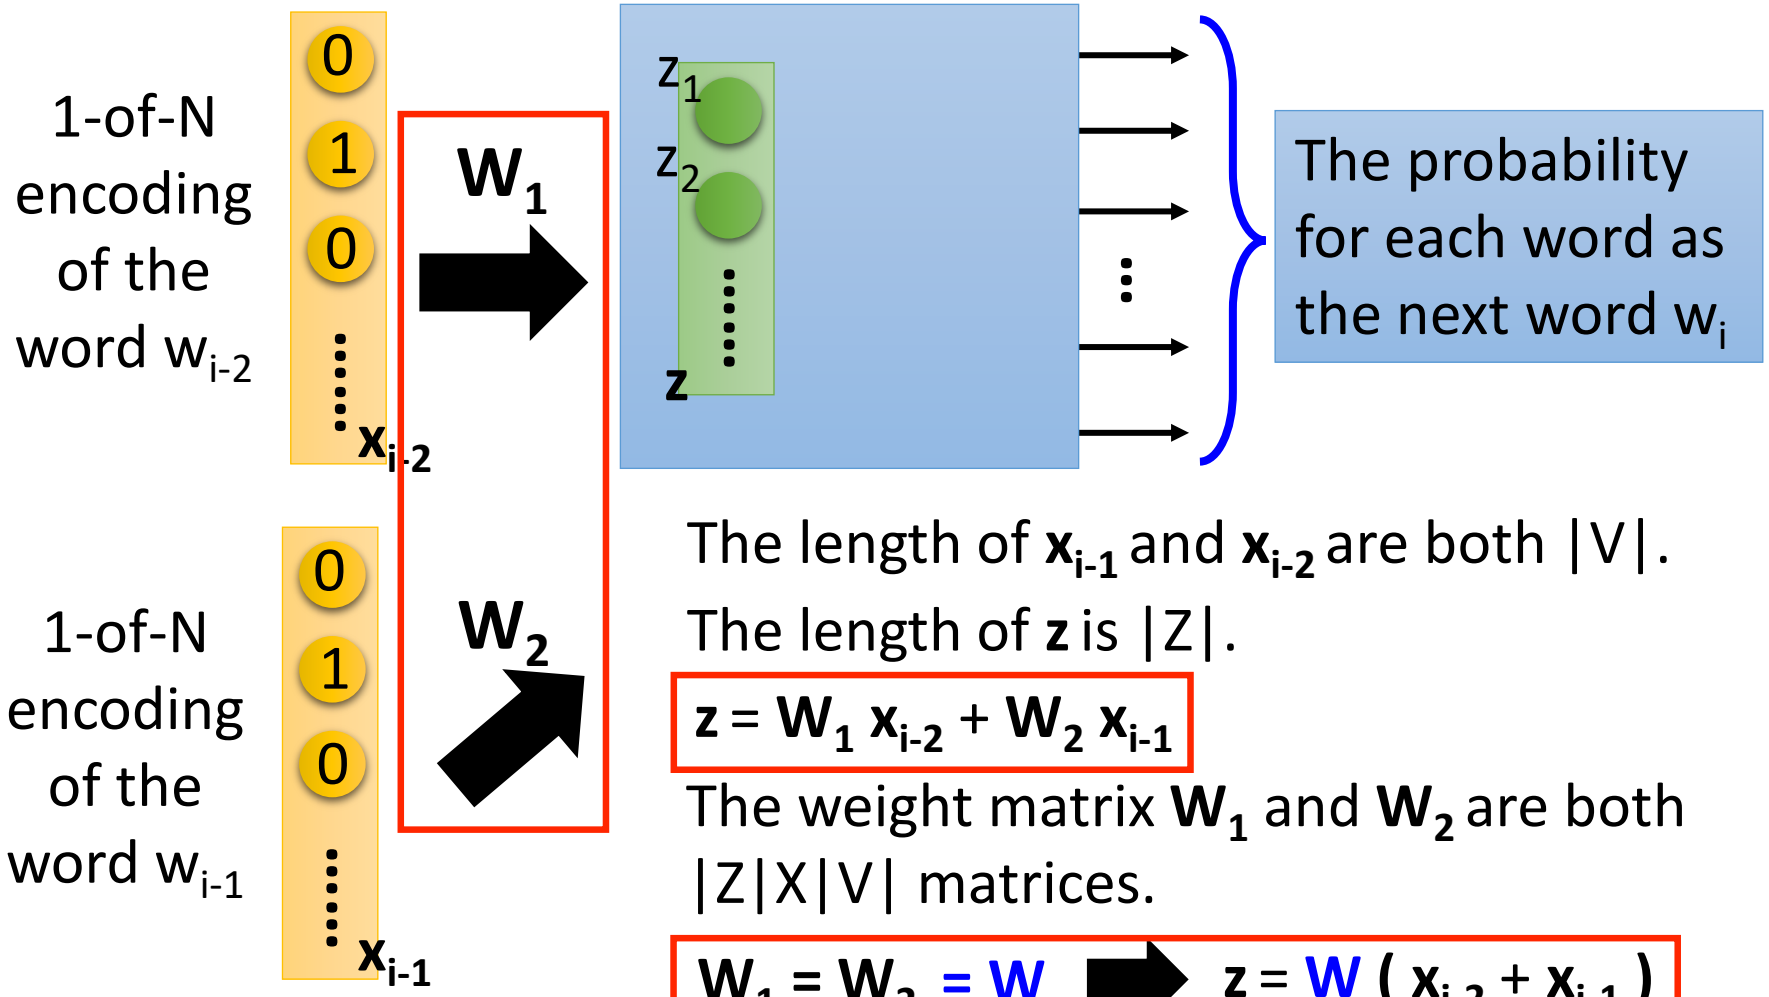

The length of $\mathbf{x}_{i-1}$ and $\mathbf{x}_{i-2}$ are both $|V|$.

The length of $\mathbf{z}$ is $|Z|$.

$$\mathbf{z} = \mathbf{W}_1 \mathbf{x}_{i-2} + \mathbf{W}_2 \mathbf{x}_{i-1}$$

The weight matrix $\mathbf{W}_1$ and $\mathbf{W}_2$ are both $|Z|X|V|$ matrices.

$$\mathbf{W}_1 = \mathbf{W}_2 = \mathbf{W} \Rightarrow \mathbf{z} = \mathbf{W}(\mathbf{x}_{i-2} + \mathbf{x}_{i-1})$$

Prediction-based – Training

Collect data:

潮水 退了 就 知道 誰 …
不爽 不要 買…
公道價 八萬 一…
……

Minimizing cross entropy

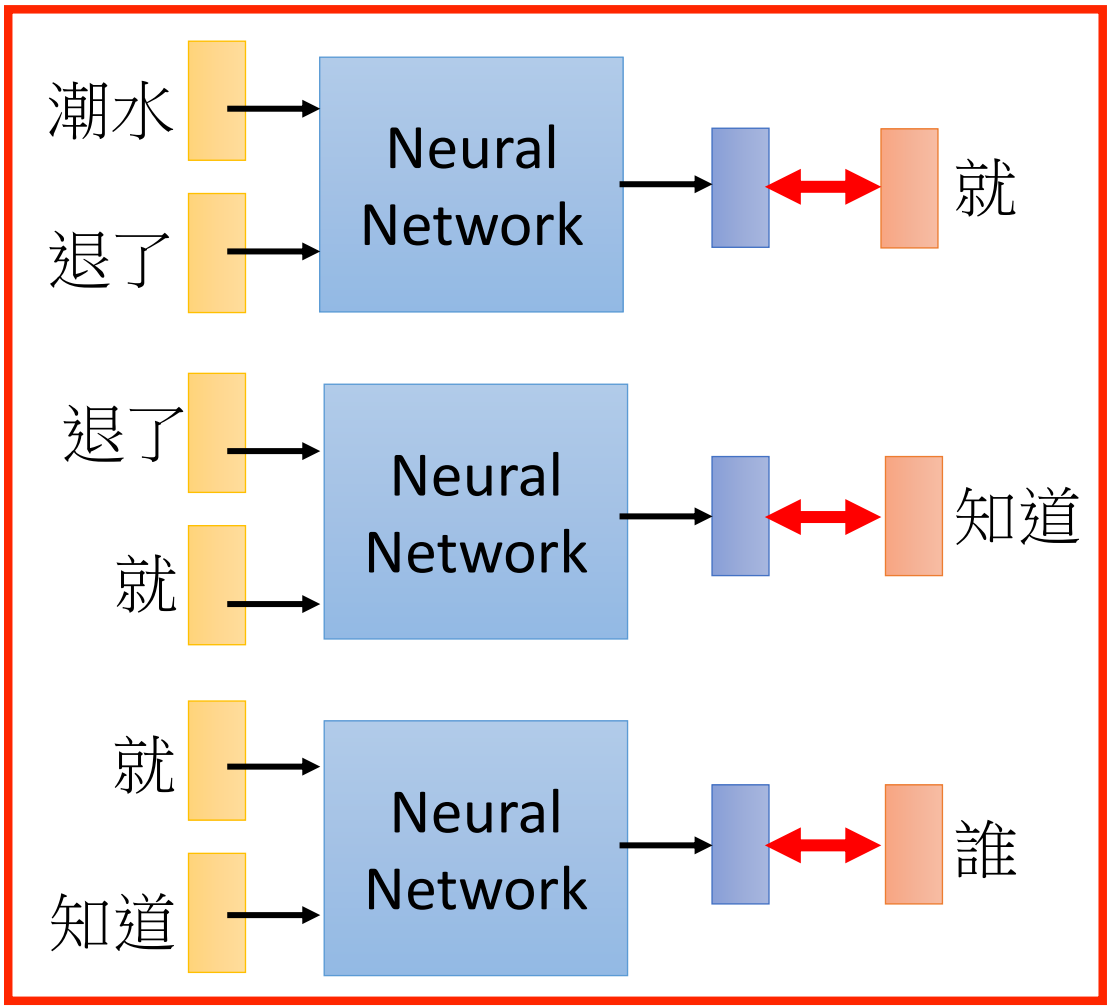

Prediction-based – Various Architectures

- Continuous bag of word (CBOW) model

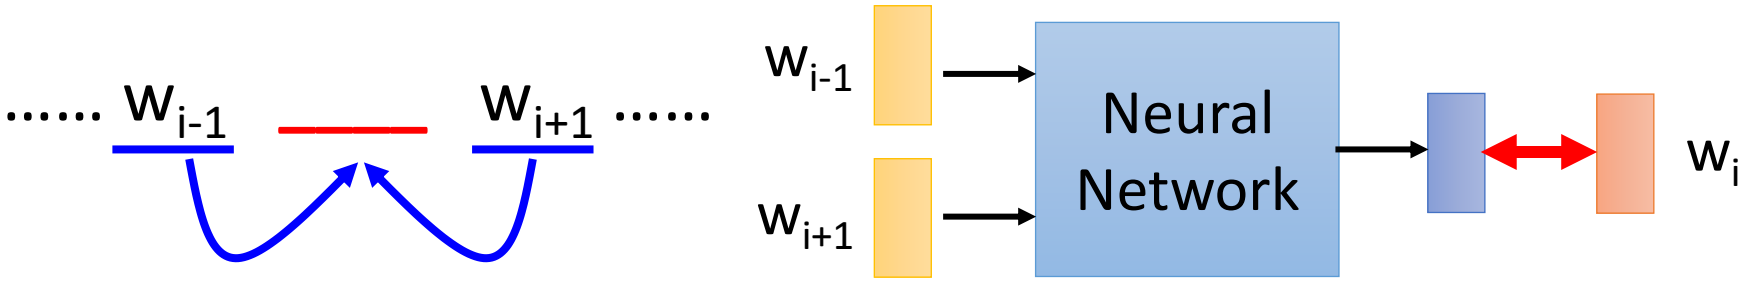

predicting the word given its context

- Skip-gram

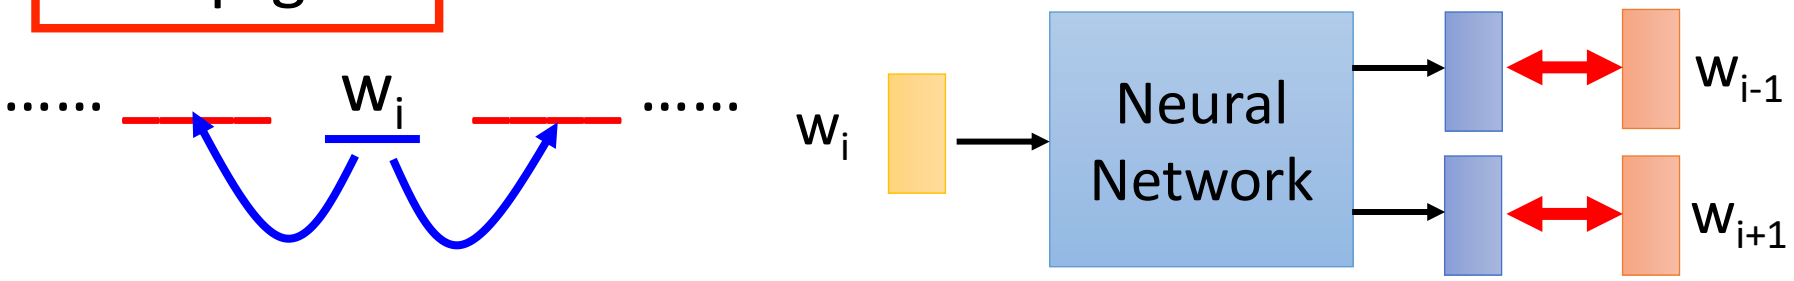

predicting the context given a word

Word Embedding

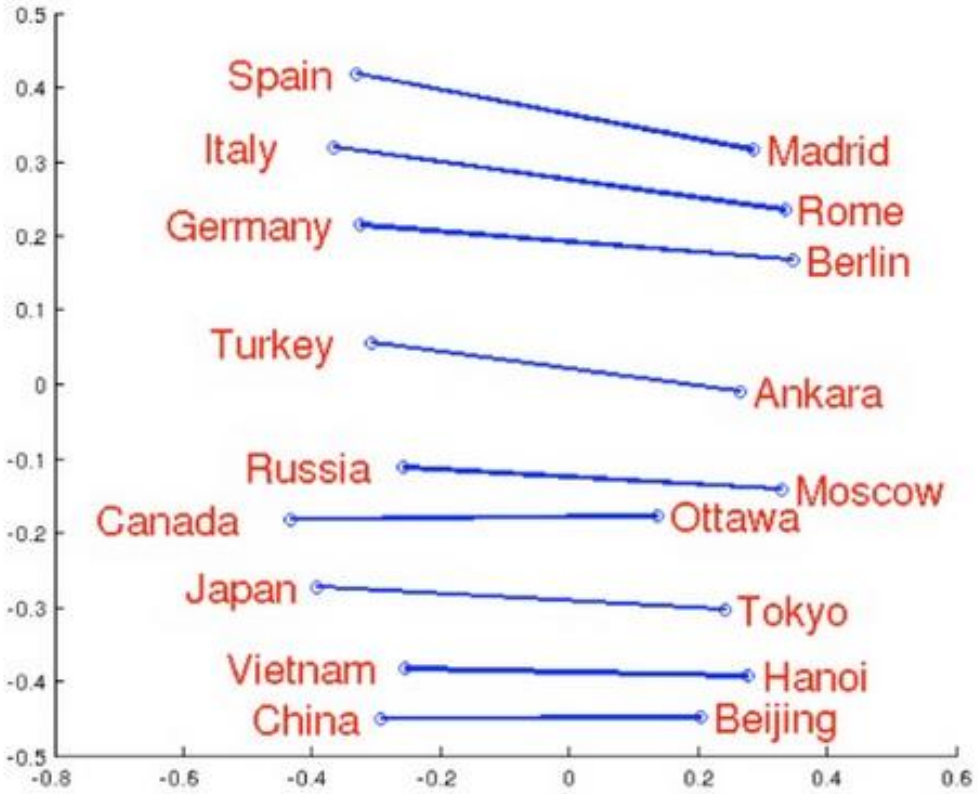

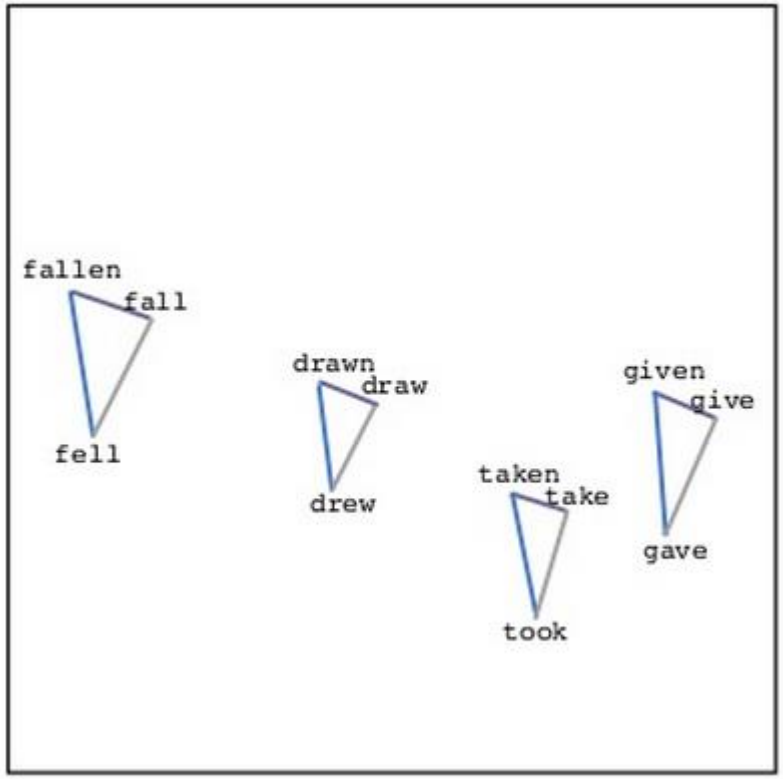

Source: http://www.slideshare.net/hustwj/cikm-keynotenov2014

Word Embedding

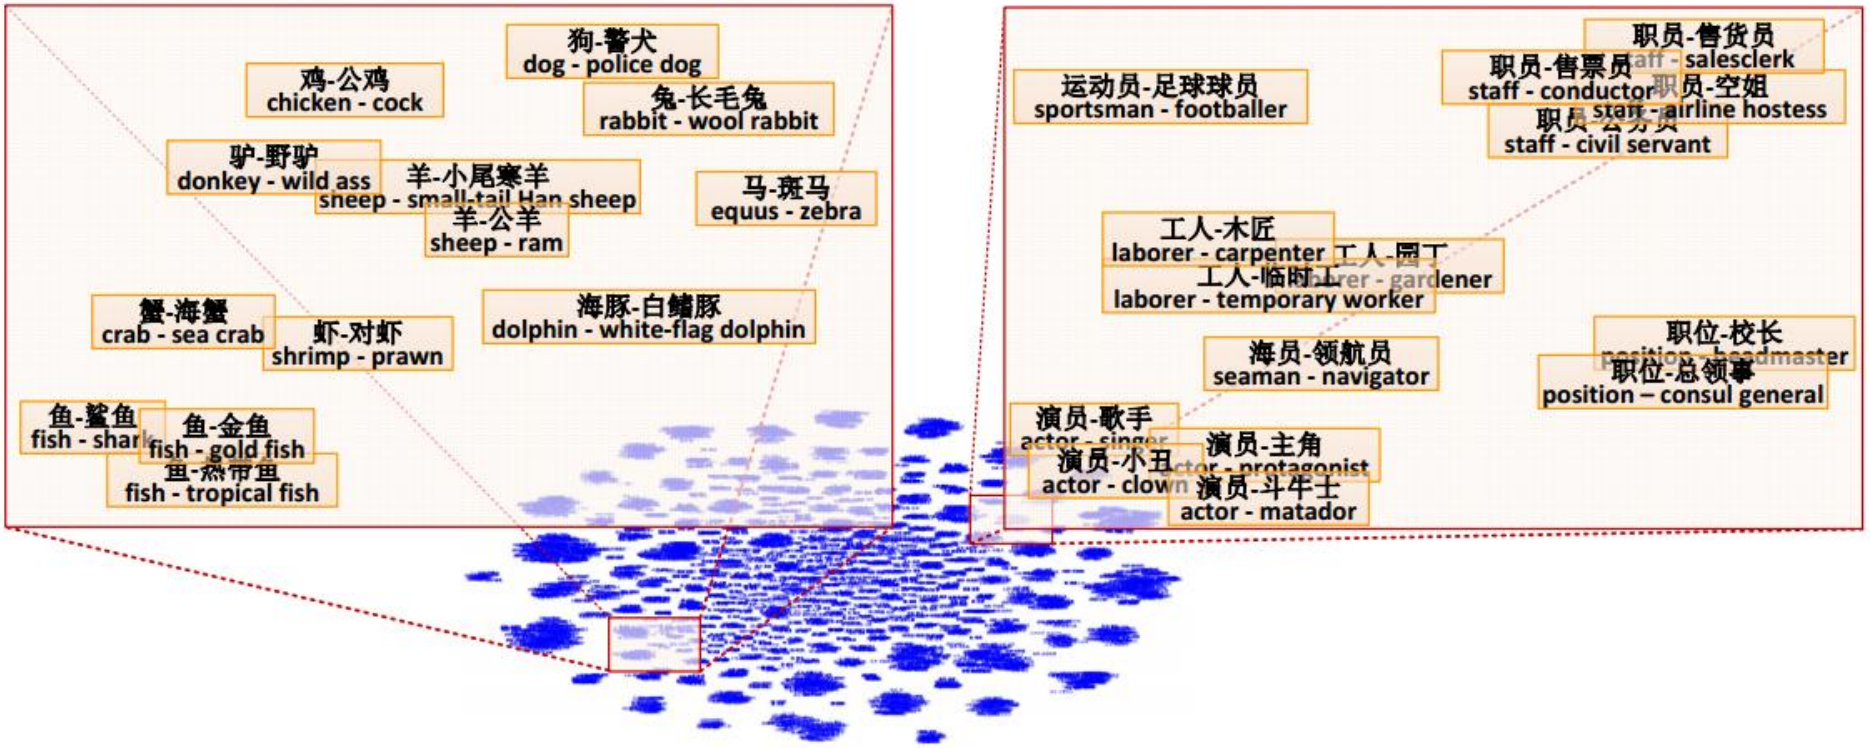

Fu, Ruiji, et al. "Learning semantic hierarchies via word embeddings."*Proceedings of the 52th Annual Meeting of the Association for Computational Linguistics: Long Papers*. Vol. 1. 2014.

Word Embedding

- Characteristics

$$V(Germany) \approx V(Berlin) - V(Rome) + V(Italy)$$

$$V(hotter) - V(hot) \approx V(bigger) - V(big)$$

$$V(Rome) - V(Italy) \approx V(Berlin) - V(Germany)$$

$$V(king) - V(queen) \approx V(uncle) - V(aunt)$$

- Solving analogies

Rome : Italy = Berlin : ?

Compute $V(Berlin) - V(Rome) + V(Italy)$

Find the word w with the closest V(w)

Demo

- Machine learns the meaning of words from reading a lot of documents without supervision

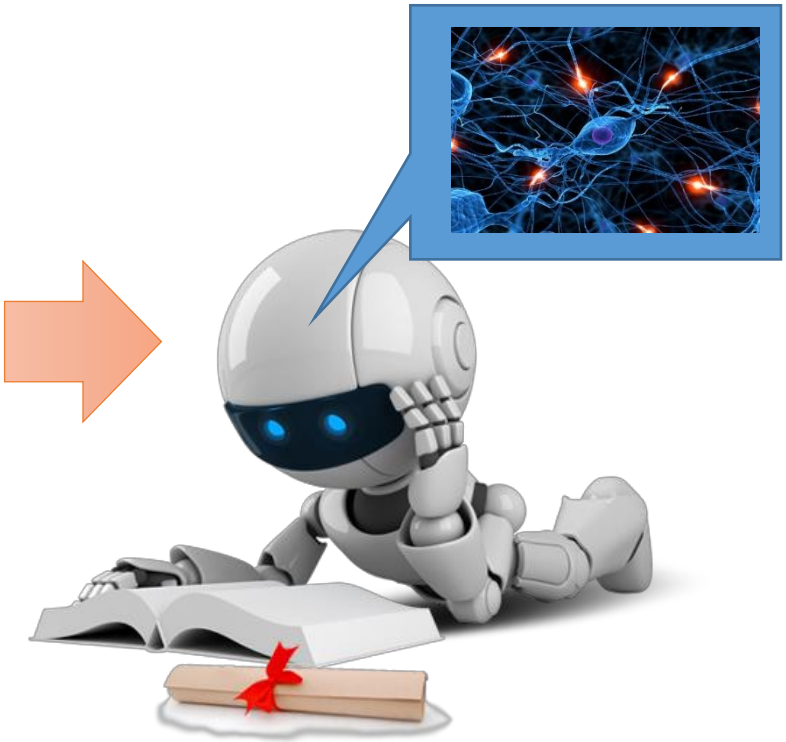

Multi-lingual Embedding

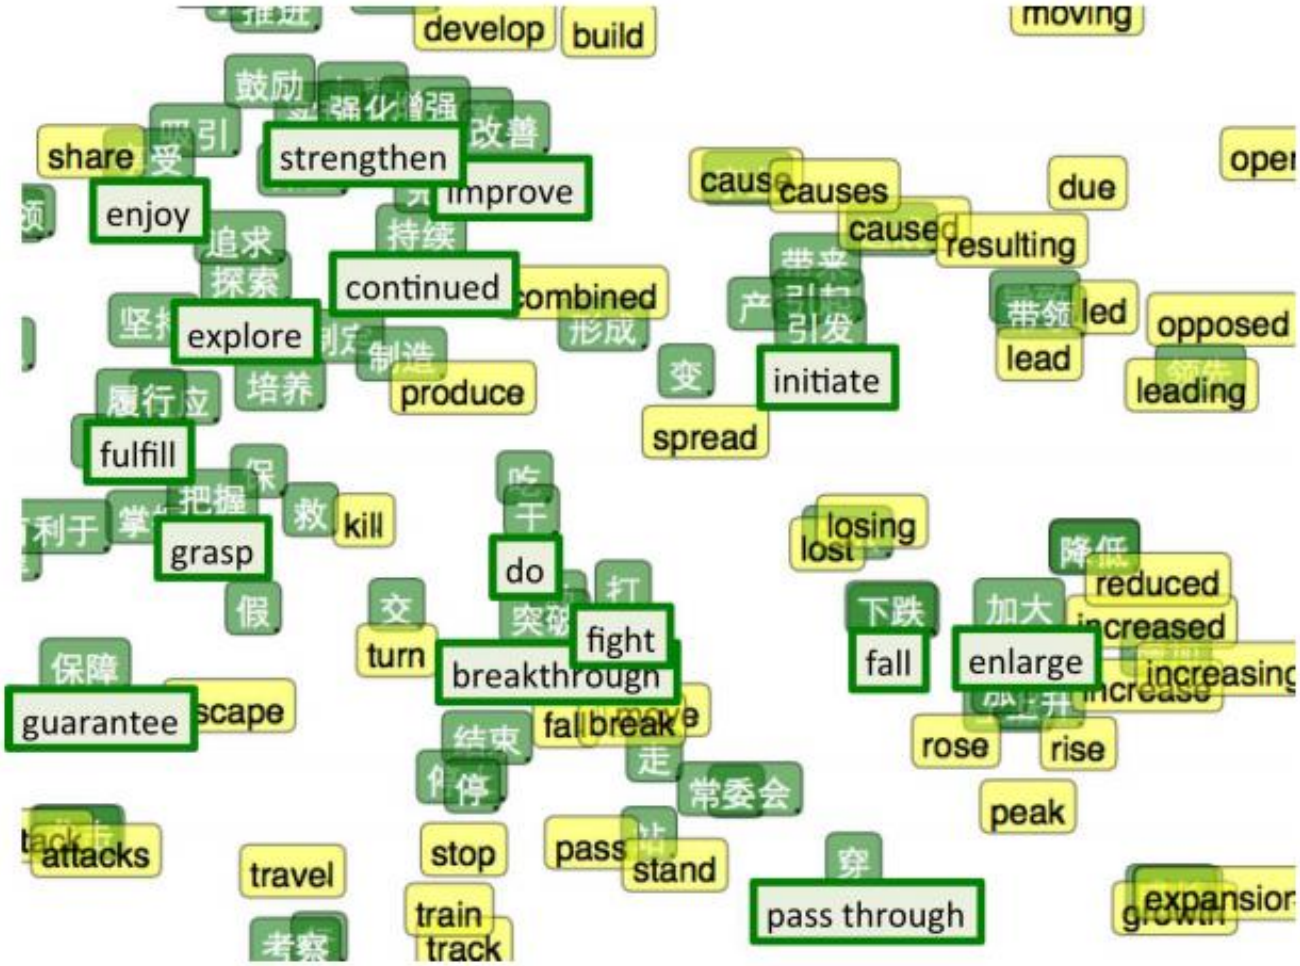

Bilingual Word Embeddings for Phrase-Based Machine Translation, Will Zou, Richard Socher, Daniel Cer and Christopher Manning, EMNLP, 2013

Document Embedding

- word sequences with different lengths → the vector with the same length
  - The vector representing the meaning of the word sequence
  - A word sequence can be a document or a paragraph

word sequence

(a document or paragraph)

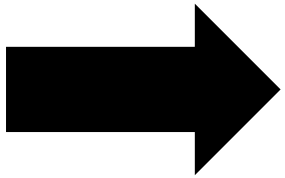

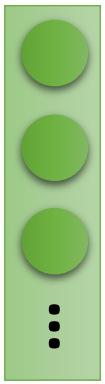

Semantic Embedding

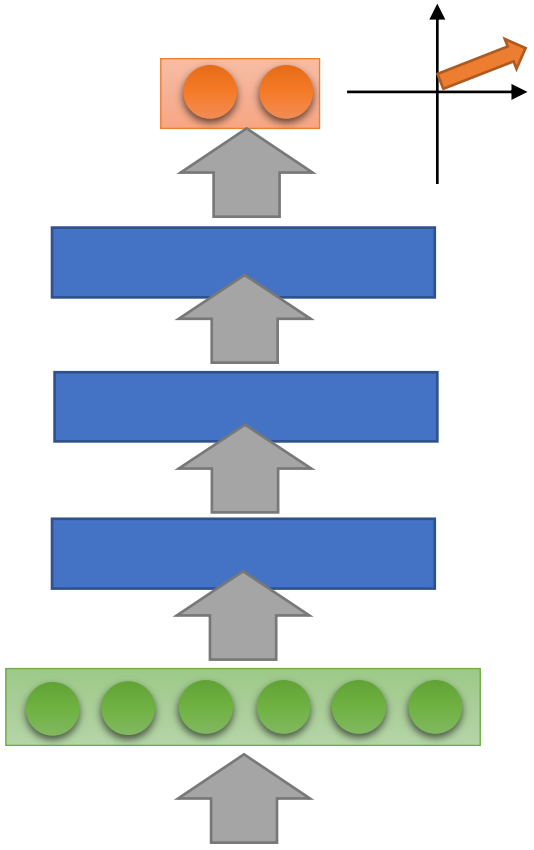

Bag-of-word

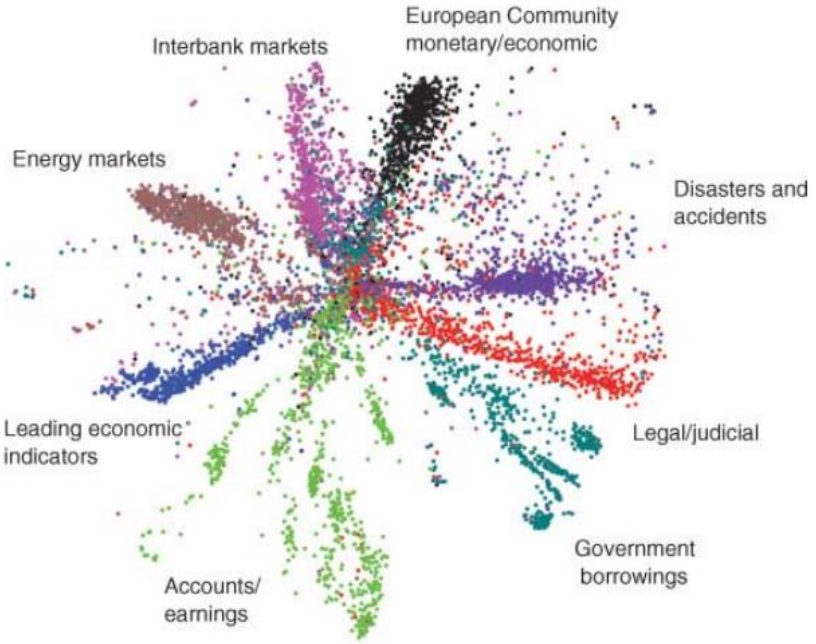

Reference: Hinton, Geoffrey E., and Ruslan R. Salakhutdinov. "Reducing the dimensionality of data with neural networks." *Science* 313.5786 (2006): 504-507

Beyond Bag of Word

- To understand the meaning of a word sequence, the order of the words can not be ignored.

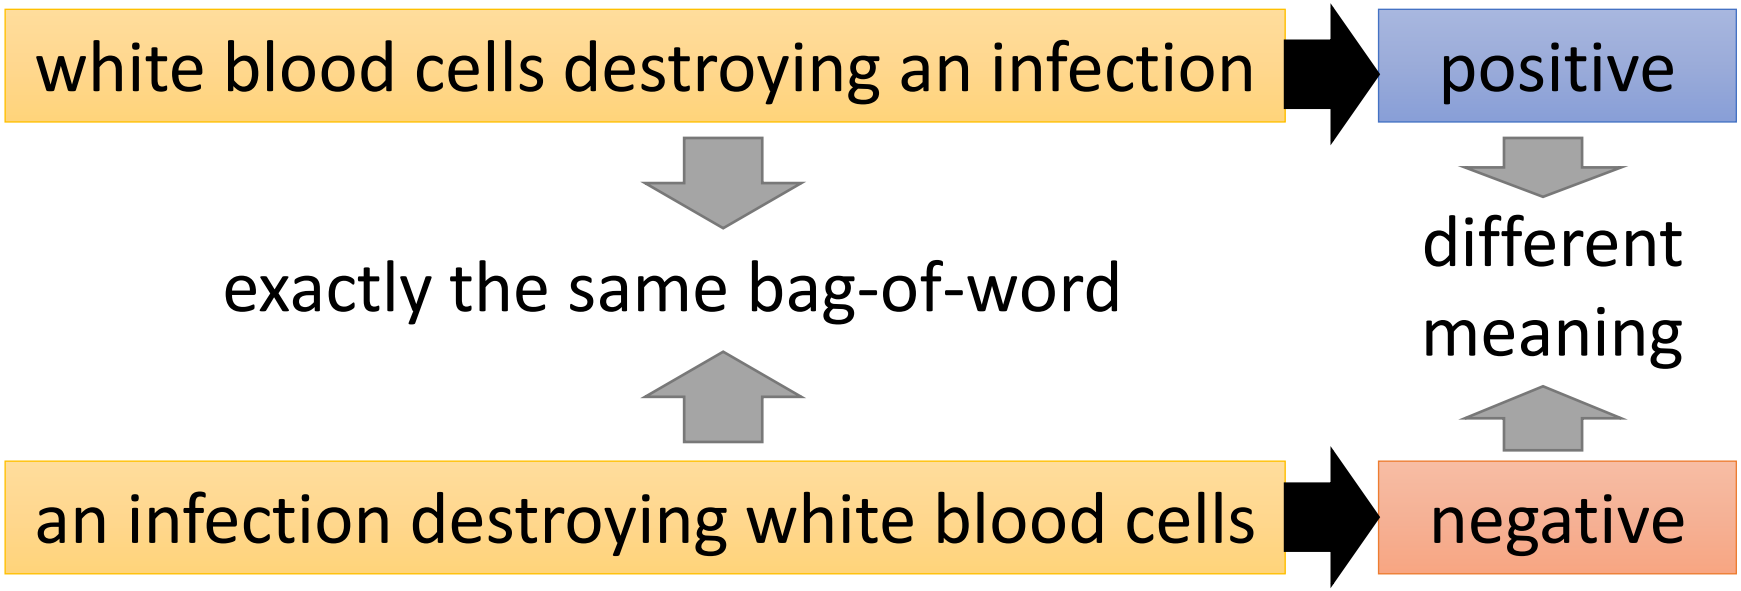

Beyond Bag of Word

- ***Paragraph Vector***: Le, Quoc, and Tomas Mikolov. "Distributed Representations of Sentences and Documents.“ ICML, 2014
- ***Seq2seq Auto-encoder***: Li, Jiwei, Minh-Thang Luong, and Dan Jurafsky. "A hierarchical neural autoencoder for paragraphs and documents." arXiv preprint, 2015
- ***Skip Thought***: Ryan Kiros, Yukun Zhu, Ruslan Salakhutdinov, Richard S. Zemel, Antonio Torralba, Raquel Urtasun, Sanja Fidler, “Skip-Thought Vectors” arXiv preprint, 2015.