

Outline

Introduction of Q-Learning

Tips of Q-Learning

Q-Learning for Continuous Actions

Continuous Actions

- Action a is a *continuous vector*

$$a = \arg \max_a Q(s, a)$$



Solution 1



Sample a set of actions: $\{a_1, a_2, \dots, a_N\}$

See which action can obtain the largest Q value


Solution 2



Using gradient ascent to solve the optimization problem.

Continuous Actions

Solution 3 Design a network to make the optimization easy.



Continuous Actions

Solution 4 Don't use Q-learning

